

Predicting employee attrition using tree-based models

Nesreen El-Rayes, Ming Fang, Michael Smith and Stephen M. Taylor

*Martin Tuchman School of Management, New Jersey Institute of Technology,
Newark, New Jersey, USA*

Predicting
employee
attrition

Abstract

Purpose – The purpose of this study is to develop tree-based binary classification models to predict the likelihood of employee attrition based on firm cultural and management attributes.

Design/methodology/approach – A data set of resumes anonymously submitted through Glassdoor's online portal is used in tandem with public company review information to fit decision tree, random forest and gradient boosted tree models to predict the probability of an employee leaving a firm during a job transition.

Findings – Random forest and decision tree methods are found to be the strongest attrition prediction models. In addition, compensation, company culture and senior management performance play a primary role in an employee's decision to leave a firm.

Practical implications – This study may be used by human resources staff to better understand factors which influence employee attrition. In addition, techniques developed in this study may be applied to company-specific data sets to construct customized attrition models.

Originality/value – This study contains several novel contributions which include exploratory studies such as industry job transition percentages, distributional comparisons between factors strongly contributing to employee attrition between those who left or stayed with the firm and the first comprehensive search over binary classification models to identify which provides the strongest predictive performance of employee attrition.

Keywords Employee attrition, Binary classification, Retention strategy

Paper type Research paper

1. Introduction

Building and maintaining a stable, productive, collaborative and high-quality workforce is a primary concern of the majority of corporate officers, as success in this area tends to be a critical contributing factor to the overall prosperity of the firm (Babington, 1993, for a survey of relevant issues). Inevitably, all firms will experience employee attrition. Involuntary attrition is often the result of profitability and performance pressures, department or business line obsolescence and mergers and acquisitions, among other factors (Datta *et al.*, 2009; Andries and Jasper, 2002; O'Shaughnessy and Flanagan, 1998). In contrast, voluntary attrition is driven predominately by employee concerns (Moninder *et al.*, 2012). Such considerations may focus around, but are not limited to, managerial direction, compensation and benefits, firm culture, firm desirability and location and promotion

The authors would like to acknowledge Andrew Chamberlain, Chief Economist at Glassdoor, for sharing the job transition data set that served as the data source for this study. The authors would also like to thank the anonymous referees and the editor Prof Peter Stokes for their comments which improved the exposition of this article.

Funding: S. Taylor would like to acknowledge partial support by the Grant Agency of the Czech Republic DyMoDiF – Dynamic Models for Digital Finance grant 19-28231X.

Competing interests: The authors have no competing interests.



potential as well as non-firm specific motivations, for example, medical conditions or retirement.

A central objective of the majority of human resource departments is to understand the root causes behind voluntary employee attrition and develop an associated mitigation strategy. Effectively navigating such issues generally results in explicit positive monetary effects stemming from increased firm revenue and cost reductions manifested through the work product of highly performant retained employees. In addition, identifying and resolving issues found to be common to employee attrition often implicitly enhances firm culture and workplace desirably, which in turn enables the recruitment of higher quality staff who further improve retention, firm operation and business practices (Cook and Ferris, 1986; Freeman, 1994). The compounding effect of the employee attrition feedback loop on overall firm success or failure provides, in our view, the essential motive to investigate the issue.

Traditionally, employee attrition- and retention-related questions tend to be examined by qualitative and anecdotal measures. Specifically, human resources staff typically conducts exit interviews after an employee provides a resignation notice to ascertain the motivations behind the decision to leave (Giacalone and Duhon, 1991). Although these conversations may be direct and candid, that is, in the event an employee is leaving for a significantly more senior role or needs to change their geographic location for family-related purposes, in actuality, human resource staff encounters considerable difficulty in discerning the employee's actual motivation. These circumstances impact the employee attrition data aggregation and quality assurance process by making it cumbersome, at a minimum, which leads to additional difficulties determining which attrition issues should be of primary importance for management to resolve. In addition, employee attrition data is generally highly confidential and only accessible to key stakeholders internally within a firm. These facts have been significant impediments to the progression of the academic literature on this topic.

There have been many recent contributions to the employee attrition literature. Specifically, Hom *et al.* (2017) survey the progression of the field from the initial study of techniques to control employee turnover to more current studies into collective employee turnover. We summarize several contributions to the employee attrition literature that are particularly relevant to the issues we investigate below. In Shaw and Delery (1998), the authors explore one of the first organizational level models of employee attrition by considering turnover patterns within 227 trucking industry organizations. Similarly, Subramony and Holtom (2012) investigate attrition patterns and associated business impact over 64 service industry firms. In Hausknecht and Trevor (2011) and Heavey *et al.* (2017), the authors organize common motivating factors for collective turnover. The authors study predictors of employee turnover in Cohen *et al.* (2015) and Fields *et al.* (2011). One of our main aims is to build upon these studies by identifying predictors of employee attrition across multiple industries and firms to understand better if commonalities exist, which may serve as focus areas for human resource professionals. In particular, we seek to broaden the set of industries considered and enlarge the sample size. We also seek to identify if we find common attrition factors such as employee salary and an employee's view of the firm's direction and culture.

Recently, internet-based platforms such as Glassdoor and LinkedIn, which are orientated toward working professionals, have amassed immense quantities of publicly available information from individual employee resumes including employment history, frank reviews of firm culture, desirability and management as well as anonymous feedback. Although this data often lacks attritional motivation information at the individual employee

level, when combined with aggregate firm culture and management rankings, one may glean many insights into the collective behavior and motivations behind individual decisions to transition to a new employer. Our major aim is precisely in this vein. More specifically, we conduct a quantitative data analysis of employee attrition motivations as well as develop predictive models that will enable human resources staff to identify employees whose firm separation may be imminent.

Several authors have studied employee attrition from a predictive modeling perspective. Early work in [Hong et al. \(2007\)](#) examined logistic and probability regression model applications to voluntary employee turnover prediction. This was extended in [Alao and Adeyemo \(2013\)](#) and [Jantan et al. \(2009\)](#), where the authors advocate for decision tree-based classification approaches and demonstrate associated improved model performance. In [Saradhi and Palshikar \(2011\)](#), the authors provide a comparison of several employee attrition classification studies on an HR Information Systems and Bureau of Labor and Statistics database. Many of these techniques were incorporated into the XGBoost ensemble model in [Punnoose and Ajit \(2016\)](#). More recently, in [Yigit and Shourabizadeh \(2017\)](#), the authors included feature engineering to improve model performance accuracy and examined tradeoffs between many traditional classification techniques including Naive Bayes, SVM, random forests, etc., on a synthetic employee churn data set provided by IBM. In more recent work, [Gao et al. \(2019\)](#) develop a weighted quadratic random forest approach to predict employee turnover at a Chinese telecommunications company. Managerial turnover is considered in [Gloor et al. \(2017\)](#), by examining patterns in email communications before resignation. Performance review appraisal is considered in [Mahmoud et al. \(2019\)](#) through decision trees and Naïve Bayes classification.

Our main aim is to continue in the direction of broadening the search over classification methods and feature engineering techniques to ultimately improve classification performance.

The main contributions of this work include an extension of [Smart and Chamberlain \(2016\)](#) where the authors examine employee attrition and retention issues based on a collection of approximately five thousand anonymously submitted resumes to Glassdoor. Specifically, we examine industry job transition patterns, dependencies between company ratings provided on the Glassdoor website and distributional-related aspects of the variables contained in this data set. We further consider how to apply modern binary classification methods to predict the likelihood of employee attrition. In particular, we examine the performance of the linear model considered in [Smart and Chamberlain \(2016\)](#), [Frierson and Si \(2018\)](#) and [Nagadevara et al. \(2008\)](#) against logistic regression, decision tree classifiers and random forests. We then extend the feature set techniques considered in [Yigit and Shourabizadeh \(2017\)](#) to include PCA-based rating features. In addition, we expand upon the comparison work in [Saradhi and Palshikar \(2011\)](#) and [Punnoose and Ajit \(2016\)](#) by performing a thorough search through a suite of dozens of binary classification methods. These models include those listed above as well as Bayes nets, SVMs with varying kernels, nearest neighbors with varying distance functions, gradient boosted trees, etc., and determine the top performers to be tree-based models from a ROC curve perspective. Finally, we delineate future data acquisition, analysis and model development extensions that we seek to investigate in future work.

The above-mentioned binary classification techniques exhibit several commonalities which we now summarize. Binary classification is the process of assigning elements of a data set into one of two possible groups based on the values of their associated features. This is achieved by initially separating the full data set into distinct training and testing subsets. Next, features defined as functions of the training data are constructed and the

classifier is fit on the training set. This fitting process varies widely according to the model under consideration. Specifically, in the case of the linear or logit regression models considered, parameters are determined by minimizing the mean squared error and maximizing the model likelihood function, respectively. In decision tree and random forest models, an initial root node is selected by splitting each feature based on a threshold and examining the resulting Gini score, which provides a measure of the purity of the resulting subset separation. The node and threshold corresponding to the minimum Gini index are selected as a root node and this process is then repeated. Random forests are constructed from an ensemble of decision trees fit on the training data set, which are then averaged to make a final prediction. Gradient-boosted trees differ by fitting an initial decision tree and then iteratively fitting additional trees to the associated residuals. Finally, there are several common performance metrics that may be utilized to compare binary classifiers. These include classification accuracy and ROC curves, which are constructed from the four possible classification outcomes. Model performance is then assessed by computing these metrics on a test set, which represents new data that was not considered during the fitting process.

This article is organized as follows: In Section 2, we describe the content of the job transition data set being considered and compile a number of summary statistics that motivate later model development. In Section 3, we pursue a more detailed examination of this data set by identifying industry transition patterns, variable importance related to attrition identification and rating variable independence. Then in Section 5, we consider several models to address the binary classification attrition problem and provide a corresponding performance comparison. Finally, in Section 6, we summarize our main findings and provide ideas for future extensions of this work.

2. Data description and summary

We first turn to describe the content extracted from a collection of employee resumes that will form the basis for subsequent attrition studies. Next, we provide a variety of summary statistics of this information that is relevant for the design of predictive models. Then we discuss our data normalization process and features constructed from this original data, which are utilized as inputs into these models.

2.1 Data source description

We worked in conjunction with the authors of [Smart and Chamberlain \(2016\)](#) to obtain a collection of 5,550 examples of employee job transitions between 2007 and 2016 which were sourced from an extensive proprietary database of resumes shared anonymously through Glassdoor's platform. A job transition is defined to be any instance of an employee listing a new role on their resume which may be associated with the current or a new employer which distinguishes between internal and external moves. Internal moves are significant in the sense that the employee was either changing roles or was being promoted within an organization. External moves are of interest to our attrition studies, as in this situation employees leave the original firm entirely.

We summarize several salient features of the data set construction process and expound upon details relevant to model development; we refer the reader to [Smart and Chamberlain \(2016\)](#) for a complete description of the data source.

Each employee job transition contains 45 attributes. Relevant attributes include employee-specific information; namely, a binary identifier indicating if the employee remained within or left their original firm, the start and end dates of employment at the original firm, the employee's average salary during their tenure with the original employer

and the employee's job title. In addition, each transition includes employer-related information. Specifically, this includes employer name and metro location, the industry sector of which the employer is a member, the founding year of the firm and the total number of employees. Finally, employer rating information is included. Particular ratings are given for the following Glassdoor created categories: overall, firm CEO, friend recommendation business outlook, career opportunities, compensation and benefits, culture and values, work-life balance and senior management performance. Finally, this data set is fully populated with the exception of missing values in approximately 6 per cent of the original and new employer founding year values; such null values are disregarded only in studies that depend upon this variable.

We note that a number of authors have considered resume data sets in human resources data analytics-related studies. For example, in [Chen et al. \(2018\)](#), the authors merge resume data sets from LinkedIn, Glassdoor.com and Monster.com. In the context of studying endorsement recommendations, [Wu et al. \(2018\)](#) manually construct a data set of approximately one thousand LinkedIn profiles. In [Luo et al. \(2018\)](#), the authors purchase a proprietary data set of approximately ten thousand resumes intended for the recruiter community.

Our data set which was constructed through the submission of anonymous resumes via Glassdoor's online portal contains similar biases. Specifically, it tends to be biased toward applicants who have recently submitted resumes, as older resumes are typically periodically removed from these platforms. Also, a disproportionate portion of the sample is oriented toward relatively high turnover positions, with a median job transition turnover of approximately 1.75 years. Finally, they do not represent a sufficient portion of the working population who have remained in the current position for a significant period of time; as a result, these type of individuals may not have uploaded a current resume to such a platform. Our data set which was constructed through the submission of anonymous resumes via Glassdoor's online portal contains similar biases.

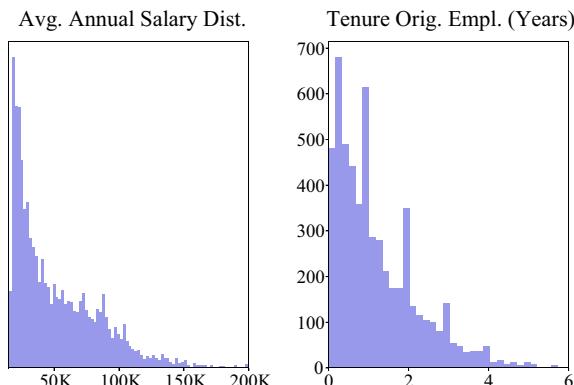
2.2 Summary statistics, feature engineering and data normalization

First, several summary statistics are presented to outline the main content of the data set that will be further explored. Next, we discuss our feature construction process to build inputs that will be important to latter exploratory studies and predictive model design. Finally, we describe the quantile normalization process that is utilized to ensure all variables are on the same scale before being input into the predictive models.

Of the 5,550 total employee job transitions in this data set, a total of 1,429 employees remained at their present firm, whereas a majority of 4,121 transitioned to a new firm. This confirms information provided in [Smart and Chamberlain \(2016\)](#), indicating that approximately three quarters of employees leave their employer during a job transition. We now graphically summarize several distributional aspects of attributes associated with these transitions.

Compensation, benefits and other forms of financial remuneration typically play a critical role in an employee's job transition decision process. In fact, an offer to substantially increase one's salary is a common impetus for a job transition. In the left subplot of [Figure 1](#), we plot the average annual salary distribution of the employee during their tenure at their original employer over our full data set. Salaries ranged from \$15,140 to \$240,000 per year. In the right subplot of [Figure 1](#), we display the tenure of each employee at their original employer before a job transition, which is similar to summary information presented in [Smart and Chamberlain \(2016\)](#).

Figure 1.
Employee average
salary and tenure
distribution at
original employer



Next, in [Table I](#), we count the industry of the original employer of all job transitions being considered and display all such industries exceeding 50 transitions. Note that the Retail and Education industries are overrepresented which provides a further indication as to why this data set includes a higher occurrence of lower salaries than the full national wage distribution. In addition, we have sufficient data to study employee industry transition patterns for many of the industries listed in this table which we explore in more detail in the subsequent section.

In [Figure 2](#), we display two additional histograms related to the original employer-specific information. In the left subplot, we present the distribution of the original firm's founding date. This histogram was left truncated to begin at 1750, with a minimum founding date of 1625 for the City of New York. Typically, firms with earlier founding dates are municipalities or governmental organizations. Note that we have an effective sampling of older and modern firms with a median founding date of 1962. In addition, we consider relatively few firms that were founded within the past 20 years in this sample as indicated by the height of the final bar of the histogram. Next, in the right subplot of [Figure 2](#), we display the log-histogram of the number of employees at each original firm being considered. The majority of employees work at larger firms that employ between ten thousand to one million people. In particular, only a small fraction of employees work at small firms with fewer than 100 colleagues. Finally, the largest employer is Walmart, with approximately 2.2 million employees.

Next, in [Figure 3](#), we display original employer violin plots of ratings data in nine categories presented for evaluation on Glassdoor's website which includes career

Table I.
Original employer
job transition counts
for industries with
more than 50
employees in our
data set

Industry	Count	Industry	Count
Retail	1,357	Manufact.	191
Education	766	Insurance	144
Info. Tech.	718	Media	113
Finance	590	Acct. and Legal	101
Bus. Services	369	Energy	92
Food Services	275	Travel	70
Telecom	248	Biotech	62
Health Care	208	Transportation	58

Predicting employee attrition

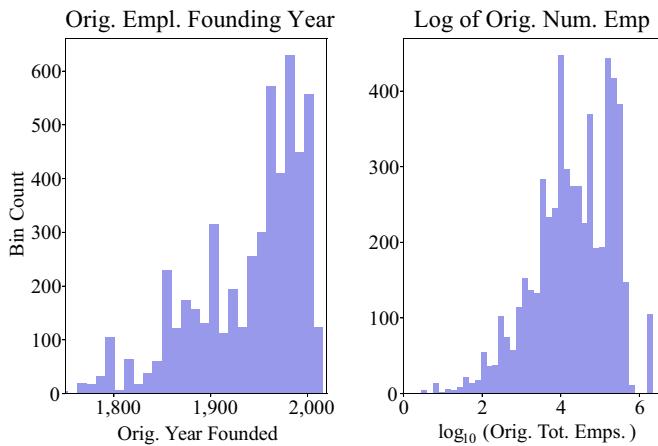


Figure 2.
Original employer
founding year and
employee number
distributions

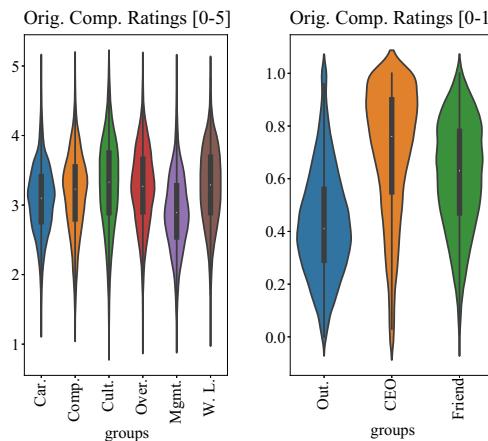


Figure 3.
Violin plots of
original employer
ratings information
sourced from
Glassdoor reviews

opportunities, compensation and benefits, company culture, overall rating, senior management, work-life balance, outlook, CEO performance and friend recommendation ratings. Here, the mean values are depicted by white circles within each form, whereas standard deviations about either side of the mean are displayed by centered black bars. The general shape of each violin is determined by a symmetric display of a kernel density estimate of the probability distribution of each rating variable. Ratings with values in the [0,5] range are plotted in the left subplot and those between [0,1] are plotted in the right. We note that no actual ratings fall outside these bounds; the slight graphical extensions beyond the boundaries in the plot are because of artifacts of the kernel density estimation procedure required to produce the visualization and are not representative of the actual data. In addition, note that one can see management ratings tend to be lower overall than other related ratings in the left subplot. In addition, cultural ratings exhibit the greatest dispersion, whereas career ratings are comparatively concentrated. The distributions of [0,1] valued ratings vary considerably. In particular, the CEO rating distribution is concentrated

to the right of the mean largely to a high occurrence of maximum ratings in approximately 9 per cent of the data. In contrast, company outlook ratings are right-skewed with a mean below the average score of 0.5. Both are less dispersed than the friend recommendation rating distribution that also is slightly orientated toward the positive side.

We now describe several elementary features that we construct from the original data that will be useful in the below exploratory and predictive studies. In particular, we will consider the percentage salary increase after a transition has been made. In addition, we will consider quantile normalized absolute changes in each rating category, for example, if an employee moved from a 75 percentile overall rating employer to an 85 percentile, we would save the 10 percentile difference as a feature. We find that these features are partially reflective of the thought process of an employee who typically leaves an organization for a higher salary and an improved company culture based on the relative rather than absolute differences in these variables and thus includes them as features. We finally note that all variables are quantile normalized in our predictive model studies so as not to bias methods because of scaling effects.

3. Exploratory insights

Now we describe a number of findings originating from the results of an exploratory analysis of the job transition data set which goes beyond the level of summary statistics. In particular, we examine to what extent salary increases motivate job transitions. In addition, if employees decide to change industry, then we study which industries they are most likely to move toward conditioned on their original industry. We then identify which variables when partitioned into employees that remained or left their firm differ the most from a distributional perspective to gain intuition on what factors should be most important for subsequent model development. Next, we investigate to what degree the nine rating categories are dependent and compute the effective dimension of these variables. Finally, we examine the relationship between firm founding year, overall original firm rating and whether or not an employee remained with or left their original firm during a job transition.

3.1 Job transition salary changes

An opportunity to earn a higher salary is often described as a primary motivation for a job transition. We seek to investigate this from a quantitative perspective and first note that approximately 13 per cent of transitions occurred without a change in salary. In [Figure 4](#), we display relative (percentage increase/decrease) and absolute salary changes. First note that this figure illustrates the need for considering features such as relative salary change, as the asymmetry of the relative change salary distribution is prominent, whereas this is not as clear in the absolute change plot. Second, note that there is a wide range of magnitudes of relative salary changes. In particular, approximately 5 per cent of employees received a salary increase of more than 150 per cent. In addition, 36 per cent took a reduction in salary as a part of their transition. We recall that our data set contains a mix of voluntary and involuntary attrition examples. Employees exhibiting involuntary attrition comprise the majority of those who encountered salary reductions.

In the extreme case, one employee transitioned from a teaching assistant to education Director at Michigan State University and received approximately a 6X salary increase. In the opposite direction, one employee transitioned from a \$230,000 salary as a Managing Director in the Education industry to a Logistics Coordinator with a \$43,600 annual salary.

Predicting employee attrition

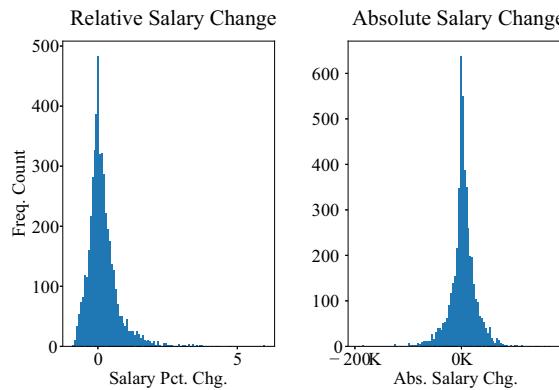


Figure 4.
Relative and absolute salary differences after a job transition

3.2 Industry transition patterns

Next, we examine how employees either choose to move to a new industry or remain in their original firm as a consequence of their job transition. In Figure 5, we display a heatmap of the percentage of employees that started in an industry indicated by the lower column labels and transitioned into the industry in the row labels. Note that both the energy and retail industries tend to retain a considerably higher proportion of their employees than the others. An interesting example along these lines is the Pharmaceutical industry of which 46 per cent of employees transition to the retail industry, while only 42 per cent remain; this is the only industry in which a higher percentage of employees transition to another industry rather than stay in the original. Moreover, the information technology industry has the lowest retention rate, with more than half of transitions out of this industry going to the financial, health care, retail and education sectors. Finally, we note that the retail industry is the most popular industry to transition into overall from a different original industry.

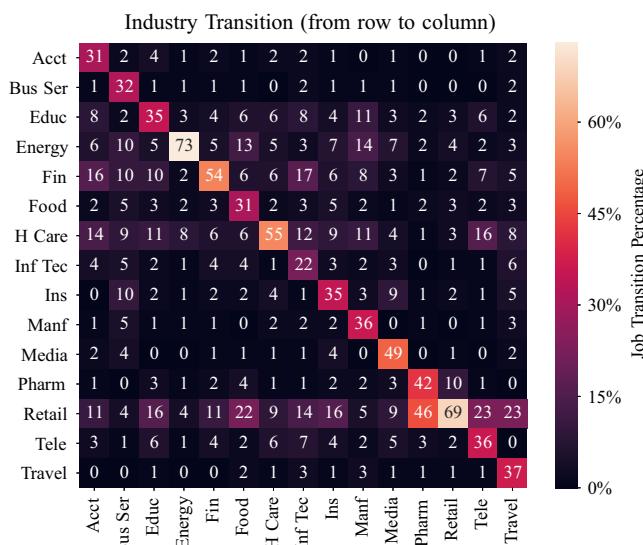


Figure 5.
Industry transition percentage: original firm industry given in columns and new/same in rows

3.3 Attrition variable importance identification

Now, we separate the job transition data set into a retention subset where employees remained with their current employer and an attrition set for those who choose to find a new employer. We compare the distributions of all numerical variables available for both groups to identify which variables have the most distinct distributions. In [Figure 6](#), we display the distributions of the original firm friend recommendation and work-life balance rating variables for employees who stayed with their original firm (green) or transitioned to a new one (red). In addition, we note that the cultural rating distributions exhibited a similar although less pronounced behavior.

3.4 Firm ratings principal components

When users fill out Glassdoor surveys that ultimately determine the company ratings provided in the job transition data set, they are asked to rate a firm in the nine categories described in Section 2.1. It is natural to allow one's overall view on the firm influence the manner in which ratings are assigned for each of these categories, that is, they are not necessarily independent. We will conduct a principal component analysis on a quantile normalized version of the full original ratings data set of nine categories to ascertain the effective dimensionality of this information.

Specifically, let r_i^j for $i = 1, \dots, n = 5,550$ and $j = 1, \dots, 9$ denote the i ratings available for j categories. Consider the sample covariance estimator of the quantile normalized ratings data

$$\hat{\Sigma} = \frac{1}{n-1} \sum_{i=1}^n (r_i^j - \bar{r}^j)(r_i^j - \bar{r}^j)^T, \quad (1)$$

where \bar{r}^j denotes the mean value of the j -th variable. Then, we perform an eigen-decomposition:

$$\hat{\Sigma} = Q\Lambda Q^{-1} \quad (2)$$

where the i -th vector of Q is the i -th eigenvector of $\hat{\Sigma}$ with corresponding eigenvalue λ_i which is the (i, i) entry of the diagonal matrix Λ . This decomposition holds, as $\hat{\Sigma}$ is assumed

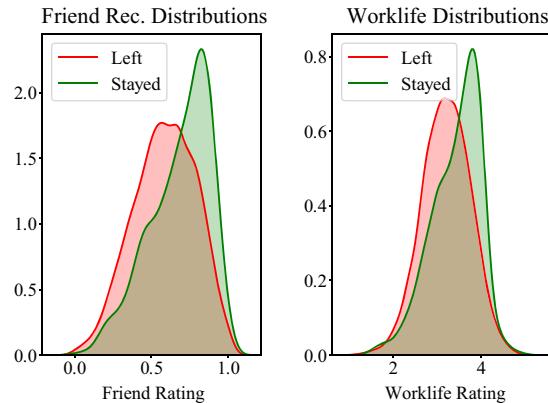


Figure 6.

Friend recommendation and work-life rating distribution for internal and external transitions

to be positive definite, and in addition, we take an ordering $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_9$ (Shlens, 2005, for an overview of principal component analysis).

In the left subplot of [Figure 7](#), we display a heatmap of the original employer ranking correlation matrix. Note that all pairs generally exhibit high correlation which demonstrates that it is necessary to account for the dependency structure of the ranking variables. In the right subplot, we display the percentage variance explained curve whose values are defined to be $w_i = \sum_{j=1}^i \lambda_j / \sum_{j=1}^9 \lambda_j$. The first principal component contains just under 70 per cent of the variation of the data and is given in normalized form by:

$$q_1 = [0.27, 0.34, 0.30, 0.28, 0.36, 0.38, 0.37, 0.37, 0.30] \quad (3)$$

which is near a uniform weighted average of the rankings with additional weight placed on career opportunities, overall rating, friend recommendation and senior management rankings. In addition, the percentage variance explained gradually increases to approximately 92 per cent at five principal components and there is no clear separation between the signal and noise portions of this data. We conclude that it will be useful to include the quantile normalized ranking weighted average in our subsequent predictive studies, but we also retain all ratings variables, given these percentage variance explained results.

3.5 Attrition, founding year and overall rating

Finally, we consider an interesting relationship between the founding year of each employee's original firm and the corresponding overall rating that they gave to the firm. In [Figure 8](#), we display two scatter plots of the original firm founding year against the overall original firm rating for the groups of employees that stayed and left their original firm in blue and orange, respectively. Note that these plots exhibit several interesting features. First, if an employee gave their original firm a high rating and the original firm year was founded before 1900, then a substantial majority of such employees remained at their original firm. In addition, as most firms with founding years in this range are municipalities,

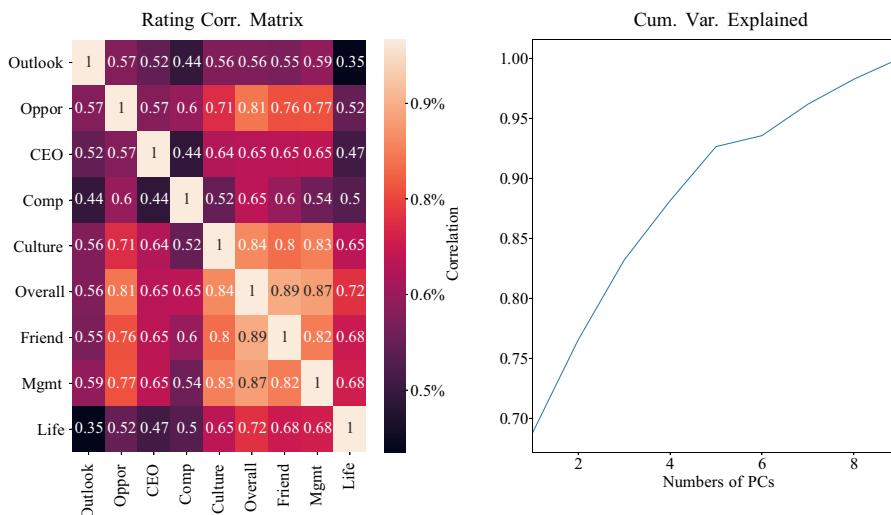
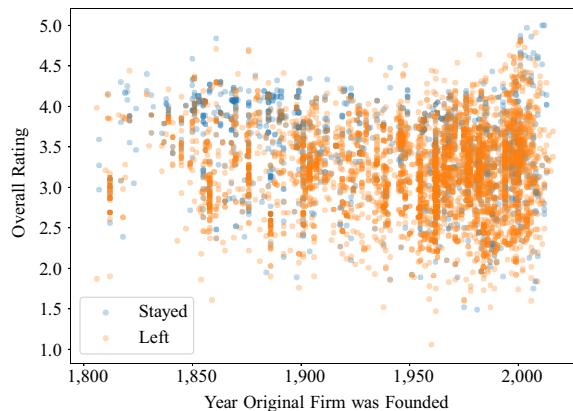


Figure 7.
Original employer
ranking correlation
matrix and
percentage variance
explained curve

Attrition vs Retention Comparison of Founding Year vs Overall Rating

**Figure 8.**

Scatterplot of overall original firm rating versus original firm founding year

this suggests government employees who are happy with their current firm are highly inclined to remain at that firm; however, we note that in the hypothesis tests, we are unable to conclude that this relationship is statistically significant because of small sample size. In addition, note that the overall ratings range varies as a function of the original firm founding year. In particular, ratings generally fall within a [2.5,4.0] interval until approximately near a founding year of 1950 where they further widen to the greatest range for newer firms founded past 2000. Finally, lower overall ratings result in higher frequencies of employee attrition than higher ratings.

4. Hypothesis formulation and testing

Based on the results of the above exploratory studies, we now formulate and test three hypotheses associated with specific causes of employee attrition. In particular, we first examine how the size of the salary increase of a job transition tends to increase the influence of the salary factor when an employee chooses to leave their current firm. Second, we seek to understand if the founding year of the firm and a high overall rating makes it more likely that an employee will remain within their current firm during a job transition. Third, we seek to understand to what extent extremely low and high overall firm ratings influence attrition. These hypotheses are stated as follows:

- H1.* Conditioned on an increase in salary, the magnitude of increase enhances the likelihood that an employee will leave their current firm despite differences between the old and new firm culture.

With this first hypothesis, we seek to understand if different levels of salary increase the likelihood that an employee will leave their current firm during a job transition. In addition, the influence of the overall rating between the old and new firm becomes less significant as relative salary change increases.

We next seek to understand if a combination of an employee's original firm founding year and that firm having a relatively high overall rating have a significant impact on attrition rates. Specifically, we test:

- H2.* Employees whose original firm has an overall rating greater than the 75th percentile that was also founded before 1900 are more likely to stay within their current firm during a job transition.

We finally examine the relationship between the magnitude of an employee's overall rating of their original firm and how it relates to attrition rates at different rating levels:

- H3.* Employees that maintain a low overall original firm rating are more likely to leave their firm upon a job transition, whereas those with higher overall ratings have a greater chance of remaining with their original employer.

Testing of the first two hypotheses will be conducted using a binary hypothesis framework with logistic regression. The main dependent variable is the relative change in the salary of the employee and we include the relative change of the firm overall rating as a control variable. We now proceed with testing *H1* by first filtering data based on a salary increase threshold, that is, job transitions that have a salary decrease less than a given percentage are removed from the sample. We refit the model by incrementally increasing this threshold and examining the *p*-values of the coefficients of both dependent variables. We note that the change in salary increase term is always significant, whereas the change in culture terms is only significant at a 5 per cent level for a salary threshold less than 40 per cent.

This result suggests that if an employee is offered a greater than 40 per cent salary increase by a competitor, then salary increase will be the predominate factor when deciding to leave one's firm.

We next test the second hypothesis by filtering the resume data set on employees who have a rating in the top quantile of all such overall ratings but also have an original employer that was founded before the year 1900. We note that the *p*-values on both variables in the analogous hypothesis demonstrate that both variables are not statistically significant at the 5 per cent level. Hence, we cannot conclude that there is a significant relationship between these two variables.

Finally, we examine the third hypothesis by binning the data into deciles based on overall original firm rating. For each bin of 555 employees, we note that assuming a fair binomial model of attrition, that according to the central limit theorem, twice the standard error of the attrition probability estimator is equal to 0.67 per cent. We display the difference between the percentage of employees who left per decile bin and the full sample mean in [Table II](#). We note that this confirms hypothesis to a 95 per cent significance level for all bins except the 61-70 bin. In particular, employees in the lowest two overall rating quantiles are 10 per cent more likely to leave their original firm upon a job transition, whereas those in the greatest quantile are 23 per cent more likely than average to remain within their current firm.

5. Toward an attrition model

We now focus on extending beyond the prior exploratory data analysis and systematically construct a series of increasingly complex but more accurate models to predict whether an

Decile bin	0-10	11-20	21-30	31-40	41-50
Excess pct	10.3	10.0	6.4	9.3	3.4
Decile bin	51-60	61-70	71-80	81-90	91-100
Excess pct	5.6	-0.6	-2.5	-14.9	-22.8

Note: Results are stated as percentages

Table II.
Excess attrition
percentage by overall
rating decile bin

employee will remain with their original employer or leave for another firm during a job transition. We consider two data sets consisting of the original variables considered in [Smart and Chamberlain \(2016\)](#) and then an extended version which includes additional features. Initially, we consider simple linear and logistic regression as well as decision tree classifiers to establish baseline results. We then more systematically search through several dozen binary classification methods to identify which have the strongest performance from an ROC curve perspective. We find that tree-based models tend to be identified in this regard and discuss the top performers in detail.

In all models considered, we uniformly randomly sample our full data set into an 80 per cent training set and 20 per cent test set. All models are trained using five-fold cross validation on the training set and all performance results are based solely on a single evaluation of each model on the hold-out test set.

We finally note that we used Python bindings of open source implementations of the classification models considered below. The majority of models may be found in the scikit-learn Python package, while privately maintained packages such as LightGBM for gradient boosted trees and tensorflow for neural net-based classification were used directly from their source made available on GitHub. In addition, we use the Python package pandas for data wrangling tasks, scipy and numpy for model implementation and matplotlib and seaborn for visualization.

5.1 Linear regression

First, we recall that in [Smart and Chamberlain \(2016\)](#), the authors fit a linear model:

$$y = \beta^T x + \epsilon, \quad (4)$$

where $\beta = (\beta_0, \dots, \beta_m)$, $x = (1, x_1, \dots, x_n)$, with random error ϵ . The target variable y is interpreted as a probability of leaving the original employer and x_i denotes normalized versions of the ratings, salary, job length and associated control variables including the employee's industry, job title and original employer's metro area.

This model is fit in the usual fashion through:

$$\hat{\beta} = (x^T x)^{-1} x^T y, \quad (5)$$

which produces a decision function for each job transition from which we may classify each employee as more likely to remain or leave their current employer; here predicted probabilities are capped at 1 and floored at 0. We include this model for purposes of comparison to [Smart and Chamberlain \(2016\)](#).

5.2 Logistic regression

We next extend to a logistic regression model with the same variables as the linear model which is more suitable for probability prediction and, in particular, binary classification applications. Specifically, we fit the model:

$$y = (1 + \exp(-\beta^T x))^{-1}, \quad (6)$$

where again the target y represents the probability of leaving the firm. Note now that $y \in [0, 1]$ which are appropriate bounds for a probability which may be considered as an advantage over linear regression-based classification. This model is fit to data through maximum likelihood estimation and we consider it as a baseline model because of its simplicity.

5.3 Decision tree classifier

Next, we explore a variety of additional binary classification models utilizing only the original employer variables described above. In particular, we examined quadratic discriminant-based classifiers, support vector classifiers and tree-based methods, among a number of other techniques. We found that decision trees provided the strongest predictive performance relative to model complexity. In particular, if we significantly increase model complexity, then we can marginally outperform a decision tree classifier; however, it is likely that such gains may not be substantive and are a result of implicit overfitting.

When constructing decision tree models, we explored trees with depths from two to ten and found that depth five trees have the strongest overall performance from an area under their associated ROC curves point of view. A decision tree is fit to the attrition data by using a greedy technique of iteratively determining which feature and threshold minimizes prediction error at each level of the tree. This method is iteratively applied until we arrive at a level of five depth tree, and then branches that only marginally contribute to the ROC curve performance are pruned.

5.4 Utilizing the full feature set

Now, we extend the previous methods by considering an extension of the prior data set. To build an analogous data set in practice, one would require access to an employee's salary increase after job transition. In addition, we add in the weighted rating PCA feature described in Section 3.4 in these studies. We note that the PCA features utilized in the predictive models were constructed only using training set data.

We consider the performance of the linear and logistic classifiers and decision tree in this case as well. In addition, we perform a broader search over dozens of binary classification methods to determine which one has the best performance from a ROC curve perspective. In particular, we now considered an ensemble of general linear models, neural network-based classifiers, support vector classifiers with a variety of kernels, nearest-neighbor classifiers and many more. The highest performing model was a light gradient boosted tree classifier with early stopping which exhibits similarities to the random forest previously considered. Specifically, gradient boosting combines many decision trees into an ensemble model in a manner that minimizes the mean squared error of predicted versus actual target values. This process is inherently iterative and the early stopping criteria terminates the process when the addition of further trees no longer improves the out of sample predictive performance of the model.

The root nodes of the trees composing this model tend to be associated with the compensation and benefits relative change feature and many branch nodes are rooted in other rating features, especially the overall rating, career opportunities and friend recommendation features.

5.5 Model performance comparison

We finally evaluate the performance of all binary classifiers considered by displaying their receiver operating characteristic curves. Each model assigns a probability of the employee remaining at or leaving the firm for each job transition. To generate the ROC curve, we set a probability threshold from which prediction probabilities below this value will be assigned to the stay group and others will be assigned to the leave group. The true positive rate defined to be the total true positive predictions divided by the sum of the true positives and false negatives is plotted against the false positive rate defined to be the number of false negatives divided by the sum of true negatives and false negatives.

First, models trained on the original variable set described in [Smart and Chamberlain \(2016\)](#) have their ROC curves displayed in [Figure 9](#). Note that the performance of the linear and logistic regression models is very similar. In addition, the five level decision tree has stronger overall performance, and the random forest which is an ensemble model of such trees has the greatest performance overall. However, the additional complexity of the random forest raises the point that the simpler single decision tree may be preferred in practice, as gains for the random forest are marginal. Here, the area under the ROC curve for the linear, decision tree and random forest is 65, 70 and 73 per cent, respectively.

We next consider models evaluated on our extended feature set; specifically, we add the rating PCA feature and the relative change in the employee's salary after a job transition. The updated ROC curves are displayed in [Figure 10](#). Here, the areas under the ROC curves are given by 67, 58, 73, 75 and 76 per cent, which demonstrates that we are able to slightly improve the performance over the prior best model in the case of the light gradient boosted tree. In addition, note that the linear regression and random forest models have very similar performance on the original data set, whereas the logistic regression model degrades in performance. The decision tree model also improves slightly. Both examples provide a strong indication that tree-based models provide a good framework for the design of employee attrition classification techniques.

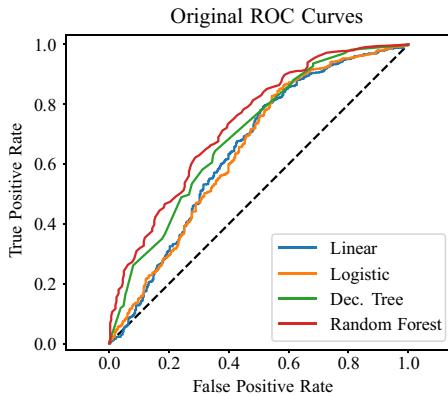


Figure 9.
ROC curves for
models trained on
original variables

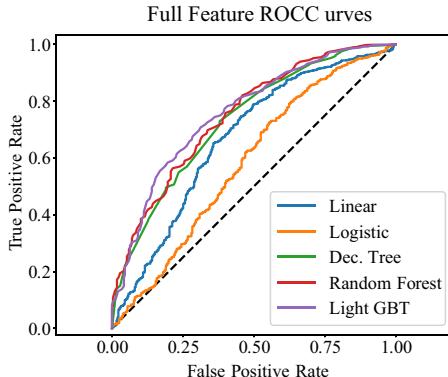


Figure 10.
ROC curves for
models trained on full
feature set

Next, we examine feature importance in the light gradient boosted tree example by randomly permuting the values of each feature individually in the training set, refitting the model and computing the area under the resulting ROC curve. We found that the average salary feature resulted in the greatest performance decline. The rating PCA feature decreased performance 31 per cent of the average salary feature. Similarly, the total number of employees and old job length had 14 and 12 per cent decreased performance when compared with the average salary feature.

Finally, we note that dozens of binary classification models were compared in this study. Of these, five strong performers include an XGBoost tree classifier, the extra trees classifier in the sci-kit learn Python package, elastic-net, l_2 regularized logistic regression and a neural-net classifier built with tensorflow; these models have AUC values of 75, 70, 69, 68 and 68 per cent.

6. Conclusions and extensions

In summary, we have obtained a data set of employee job transitions generated from anonymously submitted resumes through Glassdoor's online portal. We found several insights upon an initial study of this data which indicated that compensation, company culture and senior management performance play significant roles in influencing an employee's job transition decision. We then further investigated aspects of this data, including generating an industry job transition table, identifying which variables had the most significant changes in distribution for employees that stayed or left their current employers and constructing ratings features based on a PCA study. We then applied several binary classification models to the employee attrition problem and found that tree-based methods tended to offer the strongest performance. In particular, in the case of the original variables specified in [Smart and Chamberlain \(2016\)](#), simple decision trees offer strong performance, whereas one of their extensions, random forests, provided a marginal increase at the addition of increased complexity. Finally, we added two new features including our original PCA-based ratings feature and the percentage salary increase of the job transition.

We finally describe several ideas that we plan on pursuing in future work. First, we would like to construct a more extensive data set of employee job transition data that goes well beyond the 5,550 data points considered in this article. In particular, we hope to work with Glassdoor, LinkedIn and/or Monster.com to build a larger collection of employee job transitions and associated company ratings. In addition, we would like to obtain company-specific job transition information from LinkedIn which would permit more detailed attrition studies at the company and industry level to further the work of [Bennett et al. \(1993\)](#). Furthermore, factors such as employee engagement and absence are known to have a strong connection with employee attrition ([Kumar and Pansari, 2015](#); [Mitra et al., 1992](#)). In addition, it would be of interest to obtain firm-specific attrition data before and after attrition prevention programs were commenced. The main aim would be to assess the effectiveness of such programs and to identify if the programs that yielded the strongest results exhibit commonalities. Finally, one could then develop company-specific models, assuming sufficient data exists, that may be able to extract further information related to nuanced attrition patterns for that particular company. The results may then be merged with qualitative information gathered at exit interviews which in turn can be used by human resource staff to establish an attrition prevention plan.

We finally offer concrete recommendations to human resources professionals based on the above studies. In particular, maintaining strong Glassdoor.com ratings, in particular an overall firm rating, is crucial for employee retention. Firms in the top decile of ratings are more than 30 per cent more likely to retain employees during a job transition than those in the lowest decile. Second, ensuring that your employees are

earning a salary that is competitive with market rates is necessary for retention. If a competitor offers a significantly higher salary than the employee is presently earning, then it is highly unlikely that the employee will remain within their current firm. Finally, there is a large discrepancy between senior management and CEO Glassdoor ratings. Closer monitoring of the impact of senior management behavior especially as it relates to firm culture and impact on junior employees is advised to intervene when potentially destructive behaviors are present.

References

- Alao, D. and Adeyemo, A.B. (2013), "Analyzing employees attrition using decision trees", *Computing Information Systems and Development Informatics*, Vol. 4 No. 1, pp. 17-28.
- Andries, G. and Jasper, V.L. (2002), "The economics of skills obsolescence: a review", *Research in Labor Economics*, Vol. 21.
- Babington, P. (1993), *Building the Competitive Workforce: Investing in Human Capital for Corporate Success*, Wiley.
- Bennett, N., Blum, T.C., Long, R.G. and Roman, P.M. (1993), "A firm-level analysis of employee attrition", *Group and Organization Management*, Vol. 18 No. 4, pp. 1-19.
- Chen, L., Ma, R., Hannák, A. and Wilson, C. (2018), "Investigating the impact of gender on rank in resume search engines", *Proc. of the 2018 CHI Conf. on Human Factors in Computing Systems*. ACM.
- Cohen, G., Blake, R.S. and Goodman, D. (2015), "Does turnover intention matter? Evaluating the usefulness of turnover intention rate as a predictor of actual turnover rate", *Review of Public Personnel Administration*, Vol. 36 No. 3, pp. 240-263.
- Cook, D.S. and Ferriis, G.R. (1986), "Strategic human resource management and firm effectiveness in industries experiencing decline", *Human Resource Management*, Vol. 25 No. 3, pp. 441-458.
- Datta, D., Guthrie, J.P., Basuil, D. and Pandey, A. (2009), "Causes and effects of employee downsizing: a review and synthesis", *Journal of Management*, Vol. 36 No. 1, pp. 281-348.
- Fields, D., Dingman, M., Roman, P. and Blum, T. (2011), "Exploring predictors of alternative job changes", *Journal of Occupational and Organizational Psychology*, Vol. 78 No. 1, pp. 63-82.
- Freeman, S.J. (1994), "Organizational downsizing as convergence or reorientation: implications for human resource management", *Human Resource Management*, Vol. 33 No. 2, pp. 213-238.
- Frierson, J. and Si, D. (2018), "Who's next: evaluating attrition with machine learning algorithms and survival analysis", *International Conference on Big Data*, pages 251-259.
- Gao, X., Wen, J. and Zhang, C. (2019), "An improved random Forest algorithm for predicting employee turnover", *Mathematical Problems in Engineering*, Vol. 12(Article ID 4140707).
- Giacalone, R.A. and Duhon, D. (1991), "Assessing intended employee behavior in exit interviews", *The Journal of Psychology*, Vol. 125 No. 1, pp. 83-90.
- Gloor, P., Colladon, A., Grippa, F. and Giacomelli, G. (2017), "Forecasting managerial turnover through email based network analysis", *Computers in Human Behavior*, Vol. 71, pp. 343-352.
- Hausknecht, J. and Trevor, C.O. (2011), "Collective turnover at the group, unit, and organizational levels: evidence, issues, and implications", Cornell University, ILR School, available at: <http://digitalcommons.ilr.cornell.edu/articles/1172> (accessed 1 February 2020).
- Heavey, A., Holwerda, J.A. and Hausknecht, J. (2017), "Causes and consequences of collective turnover: a meta-analytic review", *Journal of Applied Psychology*, Vol. 98 No. 3, pp. 412-453.
- Hom, P., Lee, T., Shaw, J. and Hausknecht, J. (2017), "One hundred years of employee turnover theory and research", *Journal of Applied Psychology*, Vol. 102 No. 3, pp. 530-545.
- Hong, W., Wei, S. and Chen, Y. (2007), "A comparative test of two employee turnover prediction models", *International Journal of Management*, Vol. 24 No. 4.

- Jantan, H., Hamdan, A. and Othman, Z. (2009), "Towards applying data mining techniques for talent managements", *International Conference on Computer Engineering and Applications*, IPCSIT, Vol. 2.
- Kumar, V. and Pansari, A. (2015), "Measuring the benefits of employee engagement", *MITSloan Management Review*, Vol. 56 No. 4, pp. 67-72.
- Luo, Y., Zhang, H., Wang, Y., Wen, Y. and Zhang, X. (2018), "Resumenet: a learning-based framework for automatic resume quality assessment", *IEEE International Conference on Data Mining*.
- Mahmoud, A., Shawabkeh, T.A., Salameh, W.A. and Al Amro, I. (2019), "Performance predicting in hiring process and performance appraisal using machine learning", *10th international conference on Information and Communication Systems*.
- Mitra, A., Jenkins, G.D. and Gupta, N. (1992), "A meta-analytic review of the relationship between absence and turnover", *Journal of Applied Psychology*, Vol. 77 No. 6, pp. 879-889.
- Moninder, S., Varshney, K.R., Wang, J., Mojsilovic, A., Gill, A.R., Faur, P.I. and Ezry, R. (2012), "An analytics approach for proactively combating voluntary attrition of employees", *IEEE 12th International Conference on Data Mining Workshops*, pages 317-323.
- Nagadevara, V., Srinivasan, V. and Valk, R. (2008), "Establishing a link between employee turnover and withdrawal behaviours: application of data mining techniques", *Research and Practice in Human Resource Management*, Vol. 16 No. 2, pp. 81-99.
- O'Shaughnessy, K. and Flanagan, D. (1998), "Determinants of layoff announcements following m&as: an empirical investigation", *Strategic Management Journal*, Vol. 19 No. 10, pp. 989-999.
- Punnoose, R. and Ajit, P. (2016), "Prediction of employee turnover in organizations using machine learning algorithms", *International Journal of Advanced Research in Artificial Intelligence*, Vol. 5 No. 9.
- Saradhi, V. and Palshikar, G. (2011), "Employee turn prediction", *Expert Systems with Applications*, Vol. 28 No. 3, pp. 1999-2009.
- Shaw, J. and Delery, J. (1998), "An organization-level analysis of voluntary and involuntary turnover", *Academy of Management Journal*, pp. 511-525.
- Shlens, J. (2005), "A tutorial on principal component analysis", *Systems Neurobiology Laboratory, Salk Institute for Biological Studies*.
- Smart, M. and Chamberlain, A. (2016), "Why do workers quit? The factors that predict employee turnover", *Glassdoor Research Report Whitepaper*, pp. 1-19.
- Subramony, M. and Holtom, B.C. (2012), "The long-term influence of service employee attrition on customer outcomes and profits", *Journal of Service Research*, Vol. 15 No. 4, pp. 460-473.
- Wu, N., Xu, D. and Cho, J. (2018), "Analysis and prediction of endorsement-based skill assessment in linkedin", *2018 42nd IEEE International Conference on Computer Software and Applications*.
- Yigit, I. and Shourabizadeh, H. (2017), "An approach for predicting employee churn by using data mining", *International Artificial Intelligence and Data Processing Symposium*, pp. 1-4.

Corresponding author

Stephen M. Taylor can be contacted at: steve98654@gmail.com